

Learning Location, Not Looking: A Placement-Memorization Audit and Ten-Episode Repair at the 30M Scale

Anonymous submission

Abstract

Fixed-placement benchmarks let imitation policies learn *location* instead of *looking*. We document this on LIBERO-Object, whose target start poses are exactly pinned (measured) in 6 of 10 tasks and within $\sim \pm 1$ cm in the rest. Our platform is a 30.2M-parameter deployed stack carrying *no* language model—an open-vocabulary detector’s own text tower supplies every text embedding—making it the smallest surveyed language-conditioned policy under both size conventions at once: 17.2 M trained vs. Octo-small’s 27 M, 30.2 M deployed vs. RT-1’s 35 M even granting RT-1 the encoder its count omits. No result below is attributed to size. First, an audit finds placement memorization at three layers: a calibrated expert whose offset constant encodes pose (sign-flipping under a software rebuild), an iteration-coupled selection loop (0.700 dev vs. 0.300 held-out), and a learned grasp head whose regression rationally prefers proprioception to pixels. Second, a repair—ten teacher episodes under ± 4 cm placement teleports plus jitter on a nuisance input—restores grounding (attribution: vision 1.1–1.8 cm, proprioception 0.1 cm), verified by a substitution probe paired with a behavioural randomization protocol: the released head scores 0.400 dev, 0.700 held-out, 0.400 randomized, with both non-dev protocols confirmed at $n=50$ (35/50 [0.56, 0.81]; 26/50 [0.39, 0.65]). The protocol’s twin control cells scope that verification from both sides: on the deployment stack the memorized head scores 1/10 under the repaired head’s own shift draws (where it scores 4/10), while on a rebuilt stack whose detector shift disables the repaired head’s visual goals (0/10), the shell’s search absorbs the same shifts for the proprioception-shortcut head (6/10). The behavioural certificate is therefore joint in (head, stack). Third, the policy class that made this reachable is goal-structured decoding driving an engineered pick-and-place shell; free per-tick regression scored 0/56. All results are simulation, one task and one object in every claimed cell, $n=10$ per cell and final at that n . A dated addendum carries the multi-object campaign: four objects attempted, a second crossed and twice confirmed on never-scored seeds, two *not* crossed and —after a proposed mechanism was tested and retracted—not explained. Three pre-registered predictions were recorded and falsified, and five instruments built to quantify the boundary were all abandoned at their own calibration gates. The retracted mechanism is instructive: every object, including the two that succeed, is detected by the same generic shape prompt rather than by its name, so “the detector cannot name it” cannot be why the other two fail.

1 Introduction

A policy that scores 0.700 on a manipulation benchmark has, on the face of it, learned to see, decide, and act. This paper is about how much of such a number can instead be *location*: constants—in a benchmark’s init files, in an expert’s calibration, in a learned regressor’s shortcut—that encode where the object always is, so that vision only ever gates an approach a memorized value finishes. On LIBERO-Object [10] we measure the preconditions exactly: the target object’s start pose is bit-identical across all 50 shipped init states in 6 of 10 tasks, and varies by only 0.25–0.58 cm (std) in the other four. A benchmark with $\lesssim 1$ cm of target variance admits position-memorizers, and we found them in every layer of our own system.

The vehicle is MicroVLA, a deliberately small language-conditioned vision-language-action stack: ~ 30 M parameters deployed, with a frozen open-vocabulary detector serving as the only vision *and* text encoder. The size is the paper’s setting, not its claim. Its methodological value is that a stack this small has no capacity to paper over systems defects, so every seam had to be instrumented—and the instruments are what this paper contributes.

Contributions.

1. **A placement-memorization audit and its verification method.** A forensically worked instance of shortcut learning [6, 7] on a fixed-placement benchmark, shown at three layers, with the corrected measured account of what LIBERO-Object pins (per-task table and shipped measurement script); and the probe-plus-behavioural-randomization protocol, motivated by the on-manifold limit of substitution probes and by a counterexample to probe-only certification. The protocol’s twin control cells sharpen its scope: behavioural randomization separates memorized from repaired heads under identical draws on the deployment stack (1/10 vs. 4/10), and fails to separate them on a stack whose rebuild disables visual goals (6/10 vs. 0/10).
2. **A controlled policy-class result.** With perception, trunk, and corpus regime fixed, free per-tick regression scores 0/56 (upper 95% bound 6.4%) while goal-space supervision driving an engineered pick-and-place shell scores 4/10 dev, 7/10 held-out, 4/10 randomized—claimed strictly as a bundle, with the shell’s share measured on both stacks.
3. **Training-serving-skew instrumentation.** 29 worked defect instances at producer/consumer seams, split into disagreements catchable by parity testing and agreements-on-a-wrong-convention catchable only by provenance [3, 14].

2 Platform

Table 1: Deployed parameter ledger. “Deployed” counts every parameter loaded at inference, frozen or trained.

component	params	status
YOLO-World-S detector (vision <i>and</i> the only text encoder)	13.0 M	frozen
Tiny Recursive Model world model	9.97 M	frozen at deployment
fusion + drift + relational + planner trunk heads	7.0 M	trained (≤ 9 M cap)
goal heads (grasp 0.17 M + place 0.07 M) + learned gates	0.24 M	trained
deployed total	≈ 30 M	

No separate language model exists anywhere in the stack: the command is parsed into source/target phrases, and CLIP embeddings for (command, source, target) are harvested once per task from the detector’s own text tower. **The world model is causally inert in every nonzero number of this paper**—a persistence-TRM ablation reproduces the held-out cell exactly (0.700).

3 Results

Table 2: Protocol $\times$ head. All cells $k/10$ with Wilson 95% intervals; task 0, wrist camera. Audit-stack cells are never pooled with deployment-stack cells (§ stack finding).

protocol	memorized head	flagship (released)	sibling A	sibling B
dev (states 0–9)	7/10	4/10	7/10	7/10
held-out (states 10–19)	3/10	7/10	7/10	6/10
randomized ± 4 cm (same draws)	1/10	4/10	5/10	3/10
randomized ± 4 cm (audit stack)	6/10	0/10	—	—

Confirmed at power: held-out **35/50** = 0.700 [0.56, 0.81]; randomized **26/50** = 0.520 [0.39, 0.65]. Across a detector-stack rebuild the repaired head reproduces in all three protocols while the memorized head’s dev cell collapses 7/10 $\rightarrow$ 2/10 to its held-out floor: memorization is stack-coupled, grounding is what survived.

4 The multi-object addendum, and a falsified prediction

A second object (butter) crosses: 10/10 under an early-latch configuration, 6/10 with soup at 4/10 under a single anchor-band configuration with no per-object constants, and two independent pre-registered fresh-seed confirmations (5/10 + 3/10; 5/10 + 4/10 on a three-object head). Two further objects (cream cheese, chocolate pudding) do not cross. Goals are accurate where measurable (cream 1.58 cm on-manifold, vs. 1.34 and 1.06 for the objects that work) while each failing teacher never leaves its visual-servo phase, at source-uv std ~ 0.30 , settling tens of centimetres from the target. Four objects were attempted; two crossed.

A proposed mechanism, tested and retracted. We previously located the failure at *role binding*: the region-text head scores 0.000 on every product name these tasks are written in (measured, and correct), so groceries fall back to a shared generic tail (“box”, “cardboard box”, “can”), which would make same-shaped objects indiscriminable by construction. Screening each chain element on each object’s own corpus refutes this as a *cause*: soup resolves to “box” at 0.96, butter 1.00, cream 0.92—the **objects that cross fall back exactly as the one that fails does**, and soup succeeds 35/50 while indiscriminable in precisely the stated sense. A companion pre-registered prediction, that box-centre spread would separate them, was falsified in the same run (soup, the best object, has the *highest* spread, 0.320 vs. cream’s 0.304) and is additionally confounded by a moving wrist camera; it is abandoned, not retuned. The observation that product names never fire stands; the causal story built on it does not. The two uncrossed objects are therefore **not explained**—weaker and more accurate than the mechanism we withdrew. If every object binds through the same generic prompt, the difference must lie in the scene rather than the prompt (how many shape-alike distractors compete for it): an untested hypothesis, not a finding.

Table 3: Role-binding sub-study: four binders, three objects, $n=10$ each, one head and one configuration. Offline identity accuracy is leave-episode-out 3-way classification on corpus crops.

binder	offline identity acc.	cream	butter	soup
none	—	0/10	4/10	5/10
crop-CLIP rerank	not measured	0/10	7/10	5/10
mean prototype (centered)	0.613	0/10	2/10	8/10
1-NN bank (centered)	0.902	0/10	4/10	6/10

Offline binder accuracy dissociates from deployed success. The weaker binder owns the best soup cell and the worst butter cell; the stronger returns butter to baseline; the binder with no offline number wins butter. A matched-episode A/B probe explains it: with the 0.902 binder active the uv stream fed to the grasp head *destabilises* (std $0.258 \rightarrow 0.335$ laterally, $0.047 \rightarrow 0.202$ vertically)—a binder accurate on corpus crops thrashes on live crops. This is the paper’s own on-manifold limit recurring on two further instruments: a probe that evaluates on teacher trajectories cannot certify off-manifold behaviour, and a binder *built* from teacher trajectories cannot bind off-manifold crops.

Three predictions were pre-registered here and all three are recorded as falsified: that 0.82 per-tick identity accuracy would yield 0.914 median-of-3 latch correctness and hence cream > 0 under bank binding (the cell returned 0/10); that chocolate pudding would cross because it rarely won a misbind (its teacher never reached a grasp in five calibration iterations); and the box-centre-spread prediction above. Quantifying the deployed side of this boundary was attempted five times and abandoned five times, each at a pre-registered calibration gate: a projected-origin ground truth (rejected — soup, which succeeds 35/50, scored 0.111), its sign-corrected successor (rejected — the in-frame filter deletes the close-up target by construction), three text-tower discrimination probes (rejected — each failed to detect the working object’s own phrase), and the spread metric (rejected — confounded by camera motion, and inverted on the best-performing object). Deployed binding accuracy and text-tower discriminability are therefore reported as **unmeasured**, not estimated. The error is attributable—the prediction assumed offline per-tick accuracy equals deployed per-tick accuracy, which the probe refutes and which we do not measure. The prediction, its falsification, and this attribution are retained rather than removed.

5 Related work

Size is comparable only under a stated convention, and the two in common use disagree about which system is smallest, because published counts routinely exclude a frozen language encoder that inference nonetheless loads. We report both (Table 4). Under *trained parameters* the smallest published alternative is Octo-small’s 27 M transformer [13], and MicroVLA’s 17.2 M is $1.6\times$ below it; under *deployed* = every parameter loaded at inference it is RT-1’s 35 M [4], and MicroVLA’s 30.2 M is below that *even after granting RT-1 a free language encoder*—RT-1’s count excludes the Universal Sentence Encoder producing its instruction embedding, for which we cite no figure rather than estimate one. MicroVLA is thus smallest under both conventions simultaneously, the strongest form of the claim the evidence supports.

Table 4: Size under both conventions. † excludes a frozen language encoder the published count omits. Sizes only: capability, task breadth, and protocol are not comparable.

system	trained	deployed	language encoder
RT-2 [5]	55 B	55 B	internal (PaLI-X)
OpenVLA [9]	7 B	7 B	internal (Llama-2)
SmolVLA [16]	450 M	450 M	internal (SmolVLM2)
Octo-small [13]	<u>27 M</u>	138 M	T5-base, 111 M frozen
RT-1 [4]	35 M	<u>≥ 35 M</u> †	USE, frozen, uncounted
MicroVLA	17.2 M	30.2 M	<i>none</i> —detector’s own tower

The claim is bounded and checkable: it ranges over language-conditioned VLA policies with published parameter counts that we surveyed, it is a statement about parameter counts alone, and it is refuted by exhibiting any surveyed-class system below 17.2 M trained or 30.2 M deployed. We cannot rule out unpublished or unsurveyed smaller systems. The mechanism behind the gap is architectural rather than incidental: MicroVLA carries no separate language model at all, because the detector’s own text tower supplies every text embedding—so the frozen encoder that the other rows pay for twice (once in memory, once in the convention mismatch) does not exist here. Of MicroVLA’s 17.2 M trained parameters only 7.24 M are causally load-bearing, since the 9.97 M world model is inert in every nonzero number of this paper. No number here is comparable to community-protocol LIBERO scores (600 steps, $n=10$, wrist-only, scripted-expert corpora). Predicting *where* rather than *how* is the Transporter [18] and CLIPort [15] lineage, continued by keypoint interfaces over engineered primitives [8, 11] and, with both levels learned, $\pi_{0.5}$ [2]. Our audit is a worked manipulation instance of causal confusion and copycat shortcuts [1, 6, 17]; what we add is the audited counterexample to probe-only certification, the resulting probe-plus-randomization pairing, and the control cells that scope the pairing itself. Bootstrapping from scripted experts is established [12]; our delta is that the scripted teacher is itself audited, caught encoding placement and detector version, and repaired.

6 Claims, non-claims, limitations

Claimed: the pinning measurement in its scoped form; the released head’s cells above; free per-tick regression at 0/56; the de-memorization mechanism at the head level (attribution shift, varied-label validation, and a paired jitter ablation, $0/10 \rightarrow 7/10$ dev, exact McNemar/sign $p \approx 0.016$), scoped by the audit-stack control; 29 worked skew defects; and the bounded size result of Table 4 (smallest of the surveyed systems under both the trained and the deployed convention simultaneously). **Not claimed:** any task-success contribution from the world model (causally inert in every nonzero number); end-to-end learned motor control (the shell is engineered); head-level grounding from the randomized column alone, stack-free; object-level generalization beyond the addendum’s two crossed objects (four attempted, two not crossed and, after a tested-and-retracted mechanism, *not explained*); any deployed binding-accuracy figure; physical deployment; any size claim beyond the bounded one of Table 4 (surveyed systems, published counts, parameters only)—in particular no claim that smaller systems do not exist, and no claim that small size caused any success rate reported here. **Limitations:** $n=10$ per cell with $n=50$ confirmations on two protocols; ten held-out states reused across five disclosed selection looks; machine constants inherited from a dev-tuned era; substitution

probes are on-manifold instruments; ± 4 cm randomization sits inside the shell’s ± 6 cm probe envelope and separates heads only where visual goals are live; the teacher’s calibration consumed privileged diagnostic telemetry offline; single benchmark, single simulator.

References

- [1] Mayank Bansal, Alex Krizhevsky, and Abhijit Ogale. Chauffeurnet: Learning to drive by imitating the best and synthesizing the worst. *arXiv preprint arXiv:1812.03079*, 2018.
- [2] Kevin Black et al. $\pi_{0.5}$: A vision-language-action model with open-world generalization. *arXiv preprint arXiv:2504.16054*, 2025.
- [3] Eric Breck, Shanqing Cai, Eric Nielsen, Michael Salib, and D. Sculley. The ML test score: A rubric for ML production readiness and technical debt reduction. In *IEEE International Conference on Big Data*, 2017.
- [4] Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Joseph Dabis, Chelsea Finn, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Jasmine Hsu, et al. RT-1: Robotics transformer for real-world control at scale. *arXiv preprint arXiv:2212.06817*, 2022.
- [5] Anthony Brohan et al. RT-2: Vision-language-action models transfer web knowledge to robotic control. *arXiv preprint arXiv:2307.15818*, 2023.
- [6] Pim de Haan, Dinesh Jayaraman, and Sergey Levine. Causal confusion in imitation learning. In *NeurIPS*, 2019.
- [7] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2:665–673, 2020.
- [8] Wenlong Huang et al. VoxPoser: Composable 3D value maps for robotic manipulation with language models. In *CoRL*, 2023.
- [9] Moo Jin Kim et al. OpenVLA: An open-source vision-language-action model. *arXiv preprint arXiv:2406.09246*, 2024.
- [10] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. In *NeurIPS Datasets and Benchmarks*, 2023.
- [11] Fangchen Liu, Kuan Fang, Pieter Abbeel, and Sergey Levine. MOKA: Open-world robotic manipulation through mark-based visual prompting. In *RSS*, 2024.
- [12] Ajay Mandlekar, Soroush Nasiriany, Bowen Wen, Iretiayo Akinola, Yashraj Narang, Linxi Fan, Yuke Zhu, and Dieter Fox. MimicGen: A data generation system for scalable robot learning using human demonstrations. In *CoRL*, 2023.
- [13] Octo Model Team et al. Octo: An open-source generalist robot policy. In *RSS*, 2024.
- [14] D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, and Dan Dennison. Hidden technical debt in machine learning systems. In *NeurIPS*, 2015.
- [15] Mohit Shridhar, Lucas Manuelli, and Dieter Fox. CLIPort: What and where pathways for robotic manipulation. In *CoRL*, 2021.
- [16] Mustafa Shukor et al. SmolVLA: A vision-language-action model for affordable and efficient robotics. *arXiv preprint arXiv:2506.01844*, 2025.

- [17] Chuan Wen, Jierui Lin, Trevor Darrell, Dinesh Jayaraman, and Yang Gao. Fighting copycat agents in behavioral cloning from observation histories. In *NeurIPS*, 2020.
- [18] Andy Zeng et al. Transporter networks: Rearranging the visual world for robotic manipulation. In *CoRL*, 2020.